

Assignment -2

Abalone Age Prediction

Assignment Date	6/10/2022
Student Name	Mr.P.jeya Prathap
Student Register Number	910619104035
Maximum Marks	

1. Download the dataset

```
In [1]: # Dataset Downloaded
```

2. Load the dataset

```
In [87]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
import warnings
warnings.filterwarnings('ignore')
```

```
In [3]: df = pd.read_csv('abalone.csv')
```

```
In [4]: df.head()
```

```
Out[4]:
```

	Sex	Length	Diameter	Height	Whole weight	Shucked weight	Viscera weight	Shell weight	Rings
0	M	0.455	0.365	0.095	0.5140	0.2245	0.1010	0.150	15
1	M	0.350	0.265	0.090	0.2255	0.0995	0.0485	0.070	7
2	F	0.530	0.420	0.135	0.6770	0.2565	0.1415	0.210	9
3	M	0.440	0.365	0.125	0.5160	0.2155	0.1140	0.155	10
4	I	0.330	0.255	0.080	0.2050	0.0895	0.0395	0.055	7

```
In [5]: #Modifying the given dataset
Age=1.5*df.Rings
df["Age"]=Age
df=df.rename(columns = {'Whole weight':'Whole_weight','Shucked weight': 'Shucked_weight',
                        'Shell weight': 'Shell_weight'})
df=df.drop(columns=["Rings"],axis=1)
df.head()
```

Out[5]:

	Sex	Length	Diameter	Height	Whole_weight	Shucked_weight	Viscera_weight	Shell_weight	Age
0	M	0.455	0.365	0.095	0.5140	0.2245	0.1010	0.150	16.5
1	M	0.350	0.265	0.090	0.2255	0.0995	0.0485	0.070	8.5
2	F	0.530	0.420	0.135	0.6770	0.2565	0.1415	0.210	10.5
3	M	0.440	0.365	0.125	0.5160	0.2155	0.1140	0.155	11.5
4	I	0.330	0.255	0.080	0.2050	0.0895	0.0395	0.055	8.5

In [6]: `df.tail()`

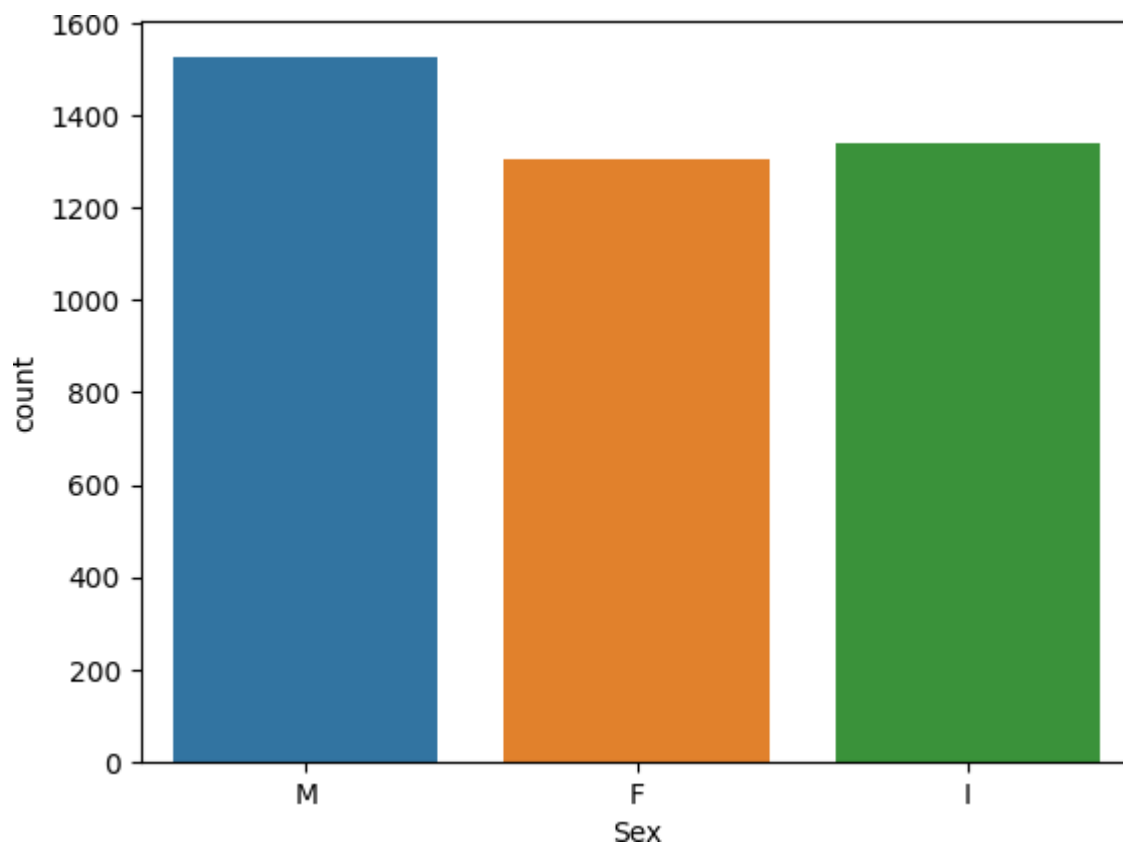
Out [6] :	Sex	Length	Diameter	Height	Whole_weight	Shucked_weight	Viscera_weight	Shell_weight	Age	
	4172	F	0.565	0.450	0.165	0.8870	0.3700	0.2390	0.2490	12.5
	4173	M	0.590	0.440	0.135	0.9660	0.4390	0.2145	0.2605	11.5
	4174	M	0.600	0.475	0.205	1.1760	0.5255	0.2875	0.3080	10.5
	4175	F	0.625	0.485	0.150	1.0945	0.5310	0.2610	0.2960	11.5
	4176	M	0.710	0.555	0.195	1.9485	0.9455	0.3765	0.4950	13.5

3. Perform Below Visualizations

Univariate Analysis

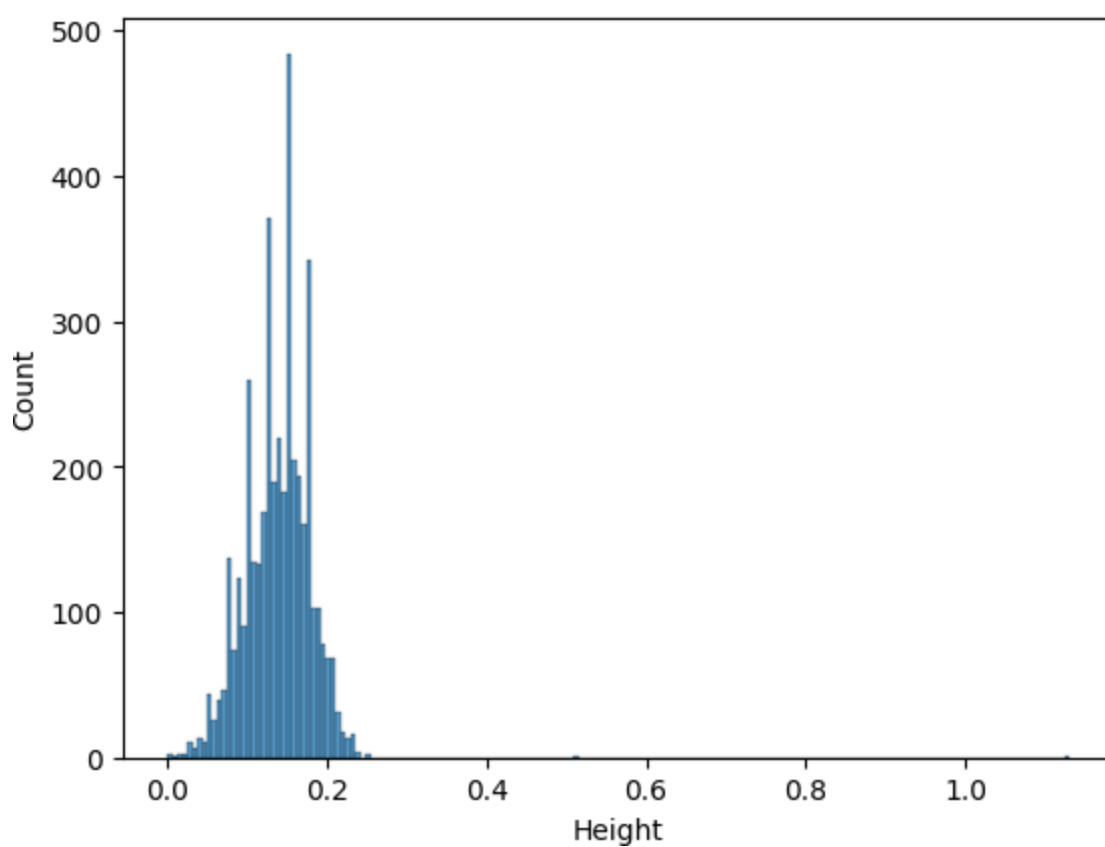
```
In [7]: sns.countplot(x='Sex', data=df)
```

```
Out[7]: <AxesSubplot: xlabel='Sex', ylabel='count'>
```



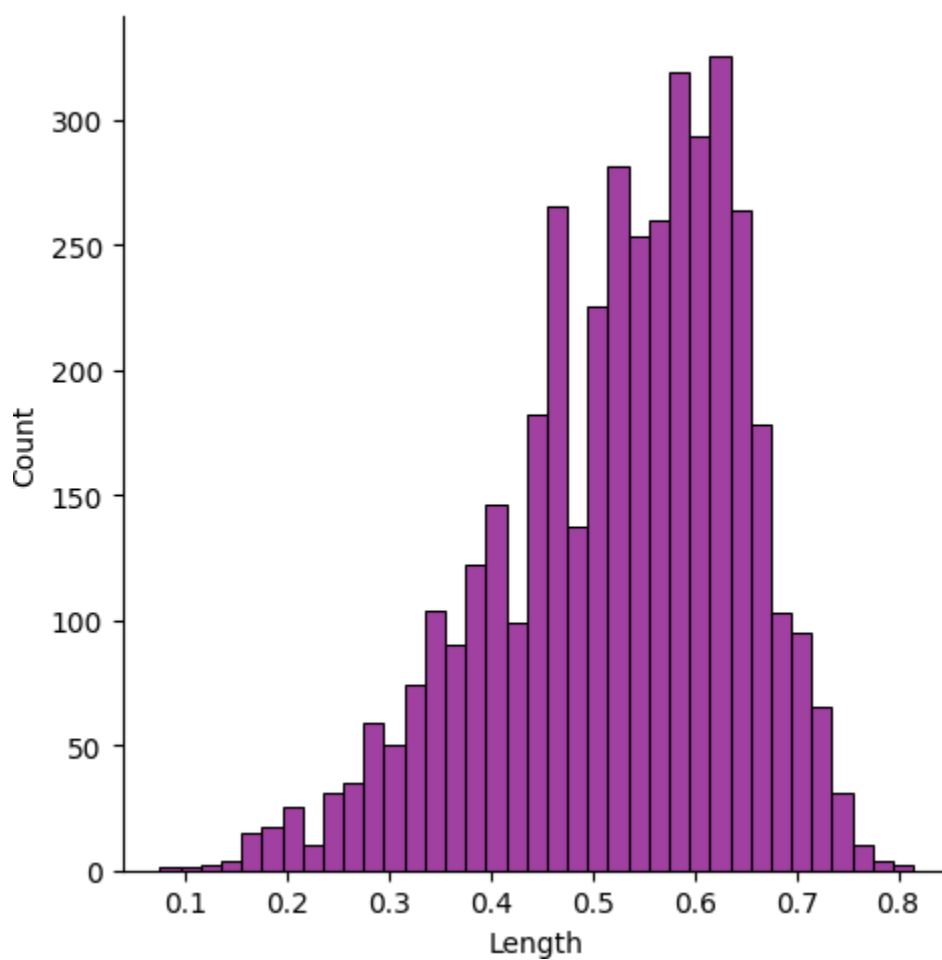
```
In [8]: sns.histplot(x='Height', data=df)
```

```
Out[8]: <AxesSubplot: xlabel='Height', ylabel='Count'>
```



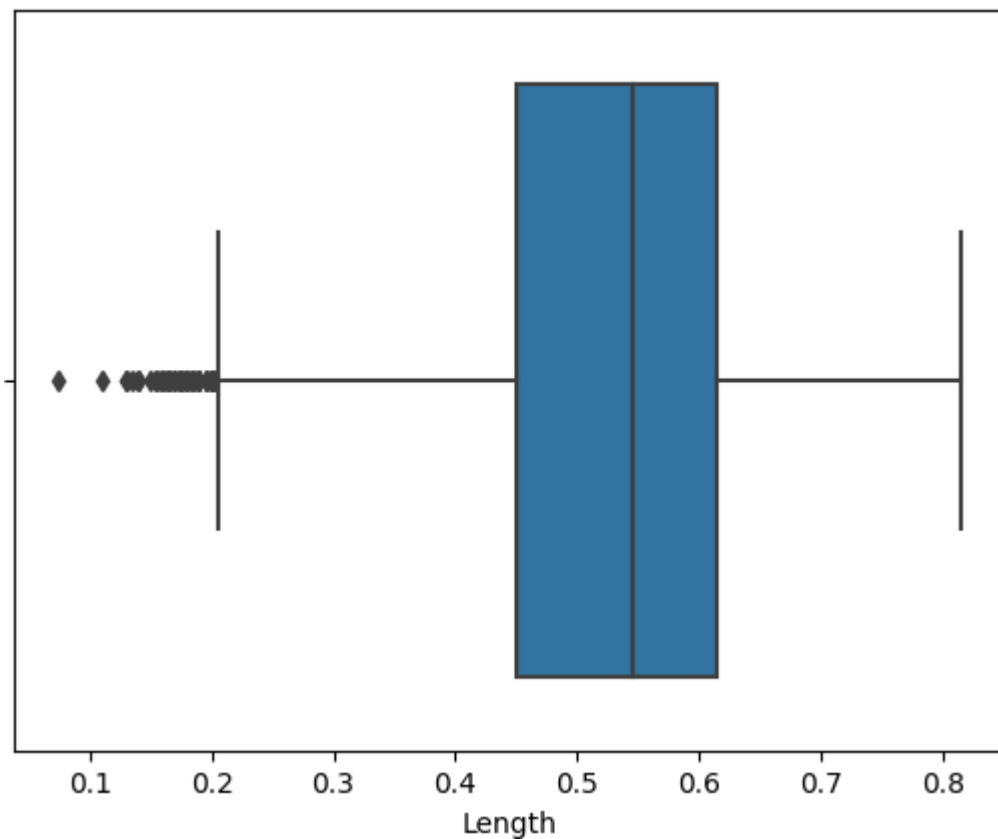
```
In [9]: sns.displot(df["Length"],color='purple')
```

```
Out[9]: <seaborn.axisgrid.FacetGrid at 0x2baa1669d50>
```



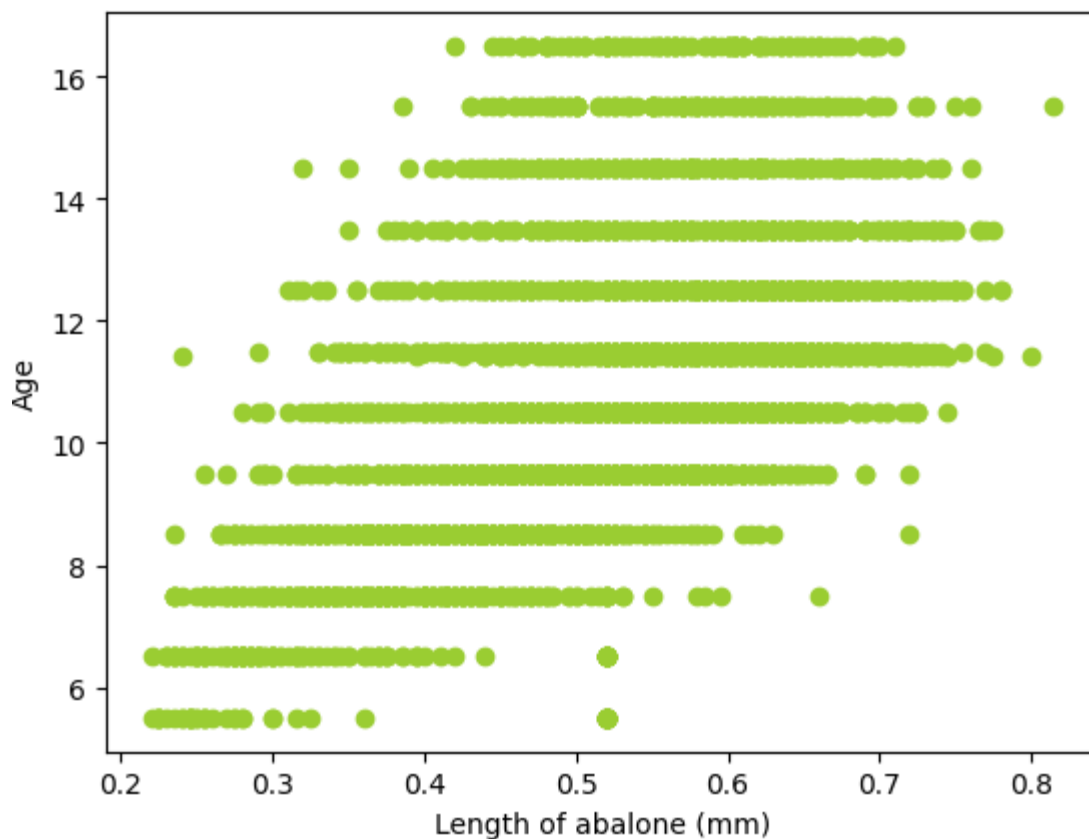
```
In [10]: sns.boxplot(x=df["Length"])
```

Out[10]: <AxesSubplot: xlabel='Length'>




Bi-Variate Analysis

```
In [69]: plt.scatter(df['Length'], df['Age'], c='yellowgreen')
plt.xlabel('Length of abalone (mm)')
plt.ylabel('Age')
plt.show()
```




```
In [70]: plt.scatter(df['Height'], df['Age'], c='crimson')
plt.xlabel('Height of abalone (mm)')
plt.ylabel('Age')
plt.show()
```

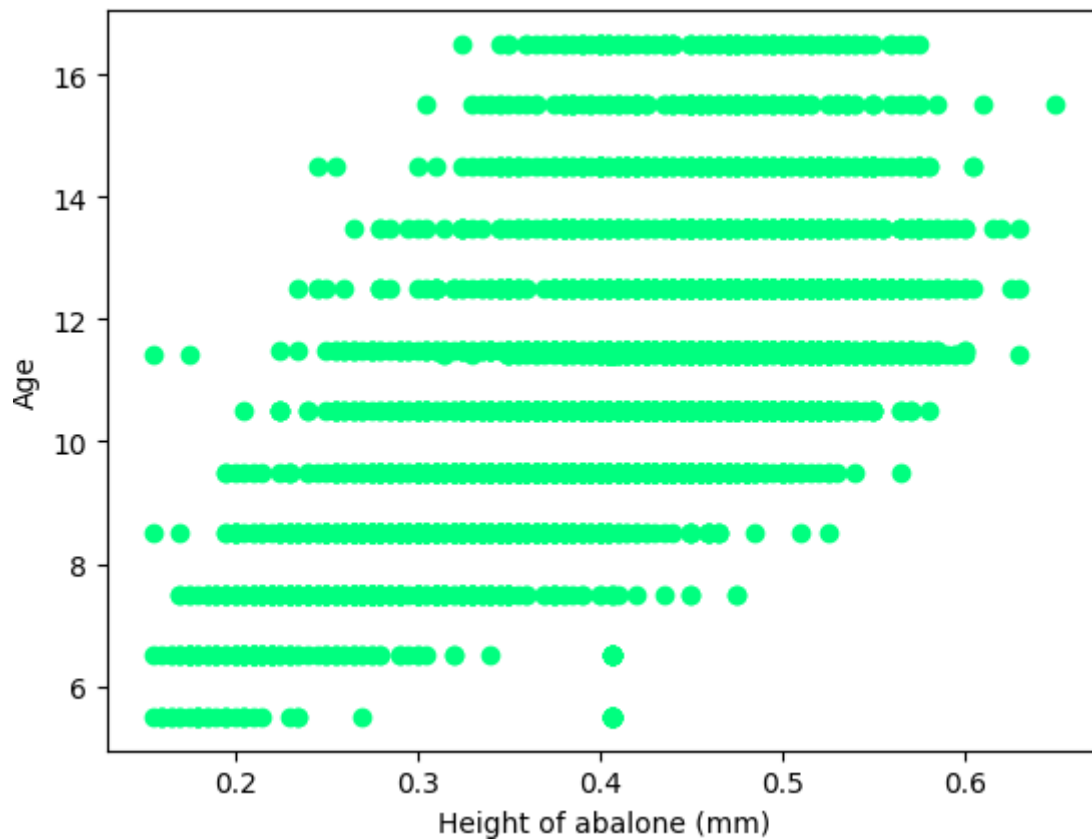


```
In [13]: sns.lineplot(x=df["Age"], y=df["Whole_weight"])
```

```
Out[13]: <AxesSubplot: xlabel='Age', ylabel='Whole_weight'>
```



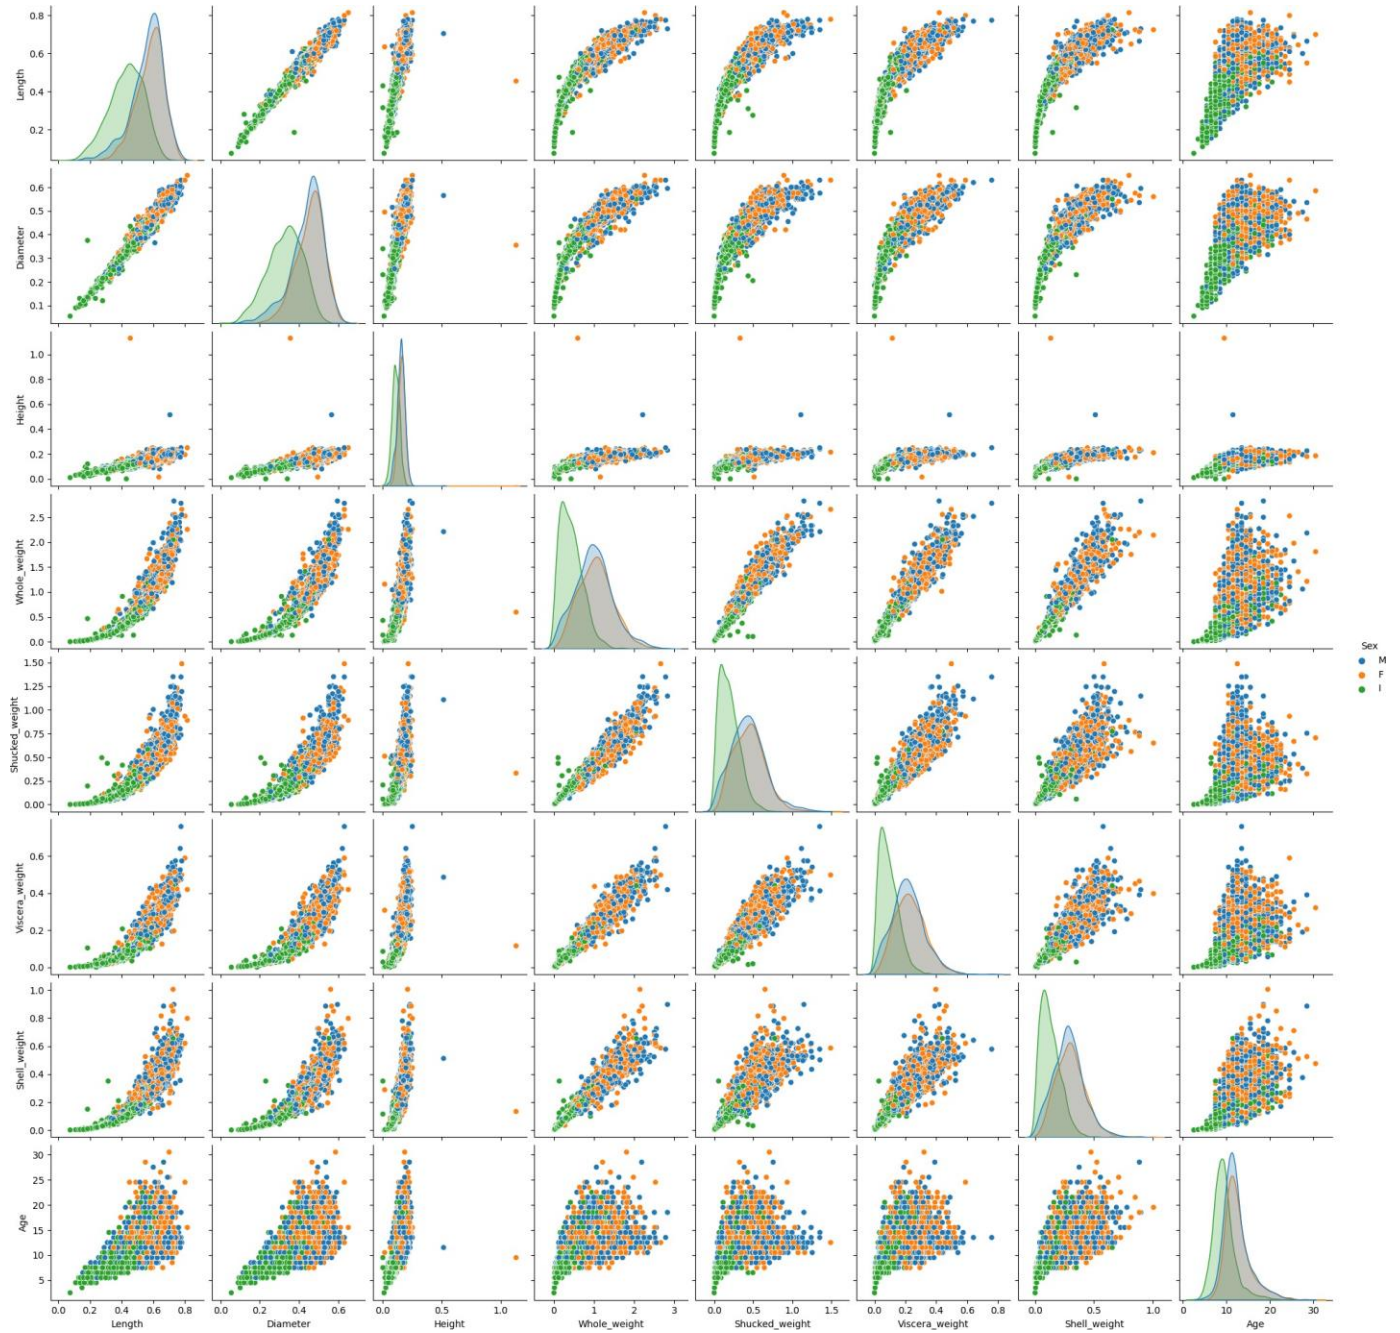
```
In [140... plt.scatter(df['Diameter'], df['Age'],c='springgreen')
plt.xlabel('Height of abalone (mm)')
plt.ylabel('Age')
plt.show()
```



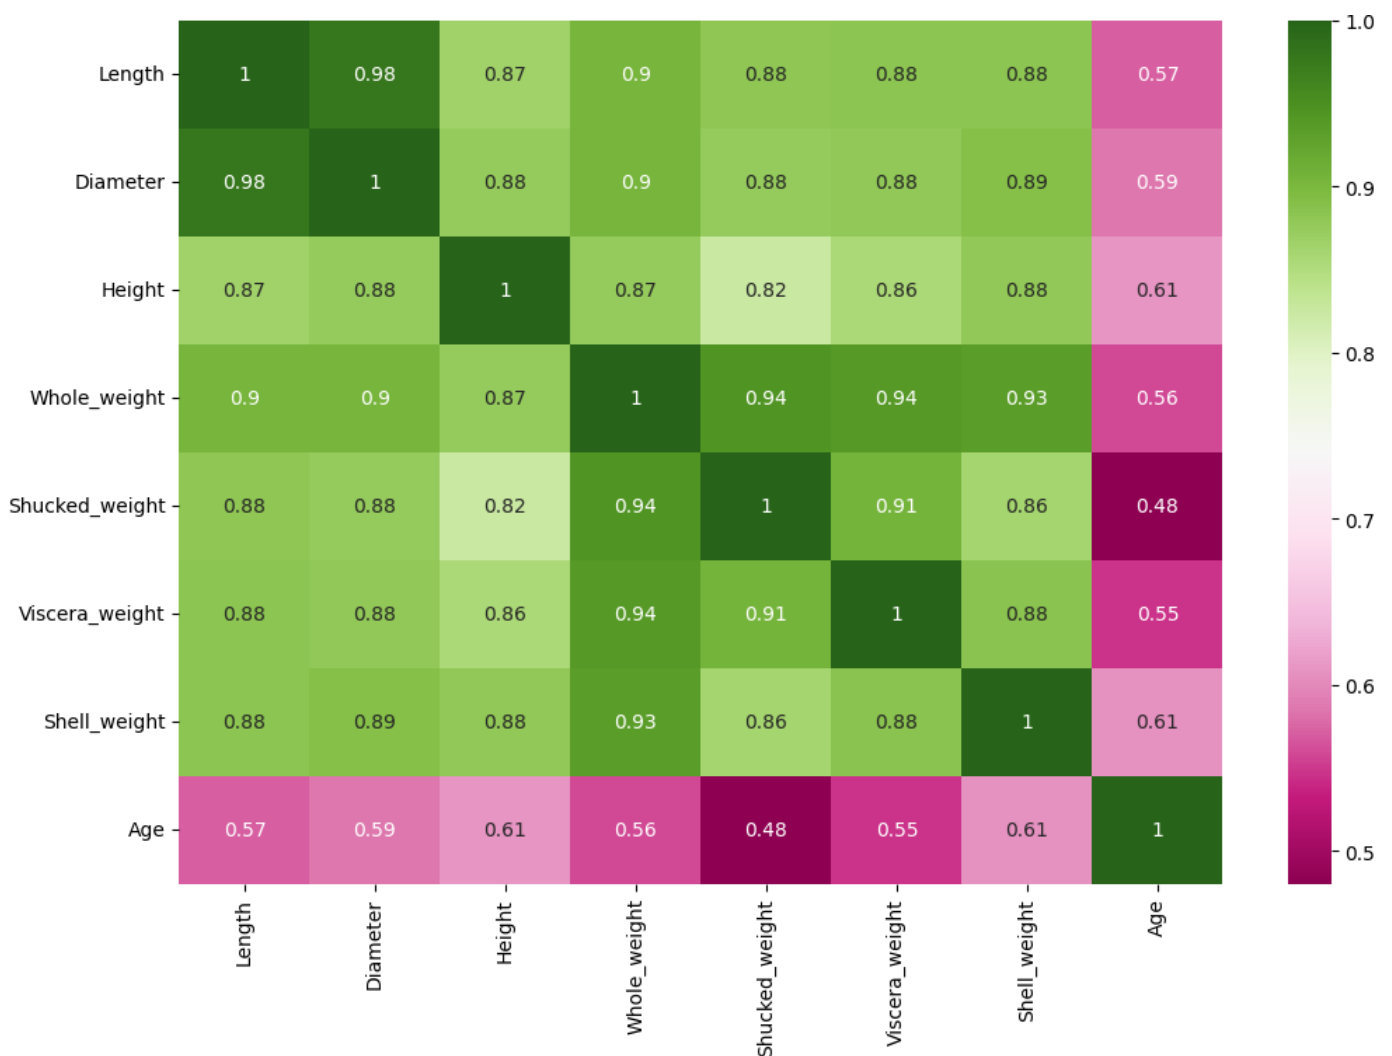
Multi-Variate Analysis

```
In [15]: sns.pairplot(df,hue='Sex')
```

```
Out[15]: <seaborn.axisgrid.PairGrid at 0x2baa2926c20>
```



```
In [88]: plt.figure(figsize=(12,8));
sns.heatmap(df.corr(), cmap="PiYG",annot=True);
```

4. Descriptive statistics

```
In [17]: df.describe()
```

```
Out[17]:
```

	Length	Diameter	Height	Whole_weight	Shucked_weight	Viscera_weight	Shell_weight	
count	4177.000000	4177.000000	4177.000000	4177.000000	4177.000000	4177.000000	4177.000000	4
mean	0.523992	0.407881	0.139516	0.828742	0.359367	0.180594	0.238831	
std	0.120093	0.099240	0.041827	0.490389	0.221963	0.109614	0.139203	
min	0.075000	0.055000	0.000000	0.002000	0.001000	0.000500	0.001500	
25%	0.450000	0.350000	0.115000	0.441500	0.186000	0.093500	0.130000	
50%	0.545000	0.425000	0.140000	0.799500	0.336000	0.171000	0.234000	
75%	0.615000	0.480000	0.165000	1.153000	0.502000	0.253000	0.329000	
max	0.815000	0.650000	1.130000	2.825500	1.488000	0.760000	1.005000	

5. Check for Missing values and deal with them

```
In [18]: df.isnull().sum()
```

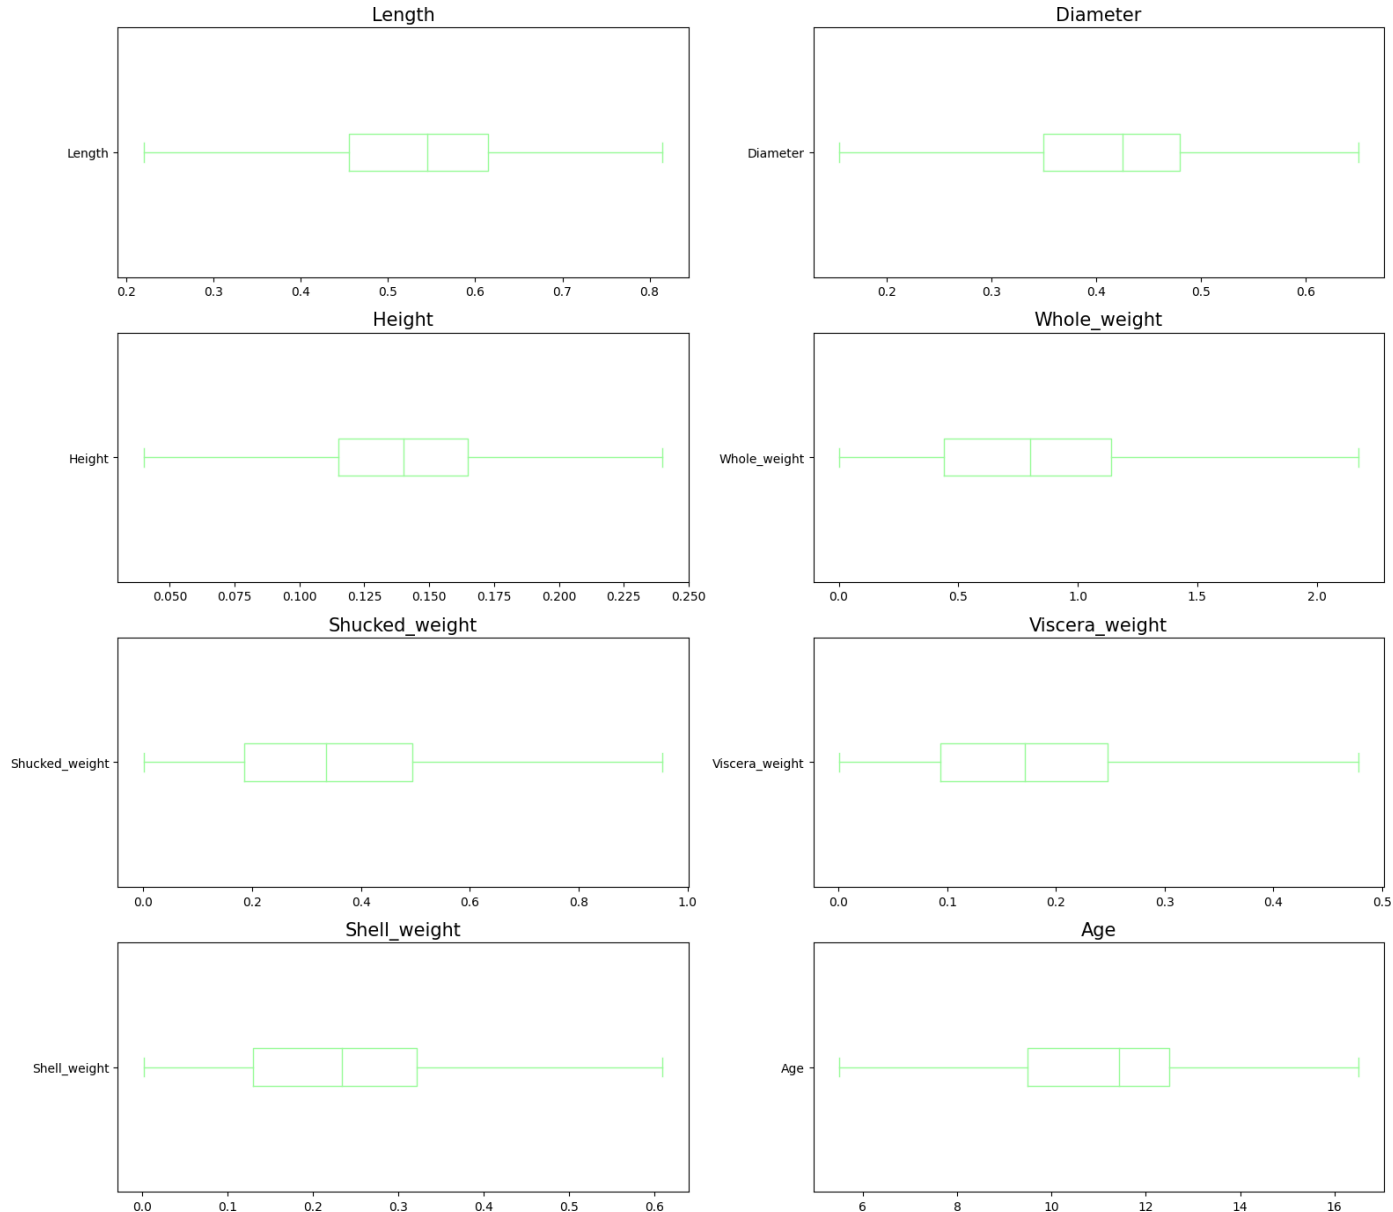
```
Out[18]: Sex      0
Length    0
Diameter  0
Height    0
Whole_weight  0
Shucked_weight  0
Viscera_weight  0
Shell_weight  0
Age        0
dtype: int64
```

6. Find the outliers and Replace their outliers

```
In [19]: col = ['Length', 'Diameter', 'Height', 'Whole_weight', 'Shucked_weight',
               'Viscera_weight', 'Shell_weight', 'Age']
```

```
In [82]: fig, axes = plt.subplots(4,2,figsize=(16, 14))
axes = np.ravel(axes)
for i, c in enumerate(col):
    hist = df[c].plot(kind = 'box', ax=axes[i],color='palegreen', vert=False)
    axes[i].set_title(c, fontsize=15)

plt.tight_layout()
plt.show()
```



```
In [89]: qnt = df.quantile([0.75,0.25])
qnt
```

```
Out[89]:
```

	Length	Diameter	Height	Whole_weight	Shucked_weight	Viscera_weight	Shell_weight	Age
0.75	0.615	0.48	0.165	1.1390	0.4935	0.2475	0.322	12.5
0.25	0.455	0.35	0.115	0.4415	0.1860	0.0935	0.130	9.5

```
In [90]: IQR = qnt.loc[0.75] - qnt.loc[0.25]
IQR
```

```
Out[90]: Length          0.1600
Diameter          0.1300
Height           0.0500
Whole_weight      0.6975
Shucked_weight    0.3075
Viscera_weight    0.1540
Shell_weight      0.1920
Age              3.0000
dtype: float64
```

```
In [91]: lower = qnt.loc[0.25] - 1.5 * IQR
lower
```

```
Out[91]: Length      0.21500
Diameter      0.15500
Height       0.04000
Whole_weight  -0.60475
Shucked_weight -0.27525
Viscera_weight -0.13750
Shell_weight  -0.15800
Age          5.00000
dtype: float64
```

```
In [92]: upper = qnt.loc[0.75] + 1.5 * IQR
upper
```

```
Out[92]: Length      0.85500
Diameter      0.67500
Height       0.24000
Whole_weight  2.18525
Shucked_weight 0.95475
Viscera_weight 0.47850
Shell_weight  0.61000
Age          17.00000
dtype: float64
```

```
In [93]: df.mean()
```

```
Out[93]: Length      0.529285
Diameter      0.411831
Height       0.139703
Whole_weight  0.815145
Shucked_weight 0.350291
Viscera_weight 0.177318
Shell_weight  0.233878
Age          10.944228
dtype: float64
```

```
In [94]: df['Length']=np.where(df['Length']<0.22,0.52,df['Length'])
```

```
In [95]: df['Diameter']=np.where(df['Diameter']<0.155,0.407,df['Diameter'])
```

```
In [96]: df['Height']=np.where(df['Height']<0.04,0.13,df['Height'])
```

```
In [97]: df['Height']=np.where(df['Height']>0.24,0.13,df['Height'])
```

```
In [98]: df['Whole_weight']=np.where(df['Whole_weight']>2.18,0.83,df['Whole_weight'])
```

```
In [99]: df['Shucked_weight']=np.where(df['Shucked_weight']>0.958,0.359367,df['Shucked_weight'])
```

```
In [100... df['Viscera_weight']=np.where(df['Viscera_weight']>0.478,0.18,df['Viscera_weight'])
```

```
In [101... df['Shell_weight']=np.where(df['Shell_weight']>0.61,0.238831,df['Shell_weight'])
```

```
In [102... df['Age']=np.where(df['Age']<5.0,11.43,df['Age'])
```

```
In [103... df['Age']=np.where(df['Age']>17.0,11.43,df['Age'])
```

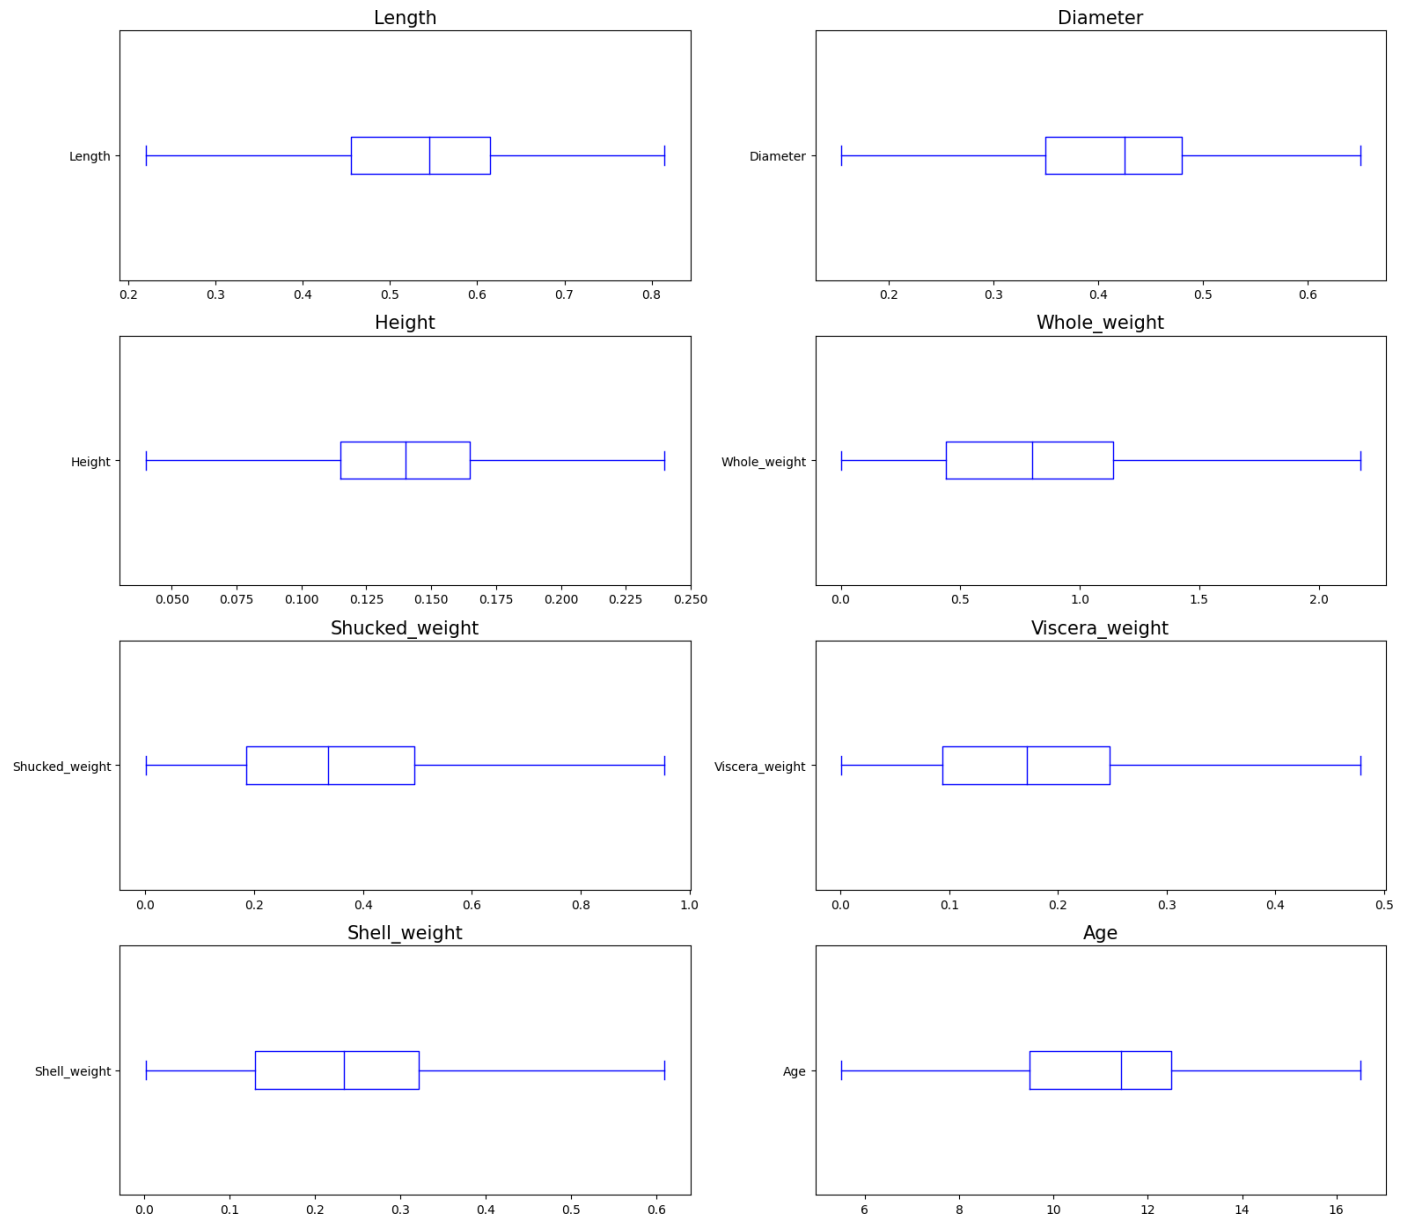
```
In [104... figfig, axes = plt.subplots(4,2,figsize=(16, 14))
axes = np.ravel(axes)
```

```

hist = df[c].plot(kind = 'box', ax=axes[i],color='blue', vert=False)
axes[i].set_title(c, fontsize=15)

plt.tight_layout()
plt.show()

```



```
In [105] df.shape
```

```
Out[105]: (4177, 9)
```

7. Check for Categorical columns and perform encoding

```
In [106] df['Sex'].unique()
```

```
Out[106]: array(['M', 'F', 'I'], dtype=object)
```

```
In [107] x = pd.get_dummies(df)
```

```
In [108] x.head()
```

Out[108]:	Length	Diameter	Height	Whole_weight	Shucked_weight	Viscera_weight	Shell_weight	Age	Sex_F	Sex_M
0	0.455	0.365	0.095	0.5140	0.2245	0.1010	0.150	16.5	0	
1	0.350	0.265	0.090	0.2255	0.0995	0.0485	0.070	8.5	0	
2	0.530	0.420	0.135	0.6770	0.2565	0.1415	0.210	10.5	1	
3	0.440	0.365	0.125	0.5160	0.2155	0.1140	0.155	11.5	0	
4	0.330	0.255	0.080	0.2050	0.0895	0.0395	0.055	8.5	0	

8. Split the data into dependent and independent variables

```
In [109... x.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 4177 entries, 0 to 4176
Data columns (total 11 columns):
#   Column          Non-Null Count  Dtype
---  -
0   Length          4177 non-null   float64
1   Diameter        4177 non-null   float64
2   Height          4177 non-null   float64
3   Whole_weight    4177 non-null   float64
4   Shucked_weight  4177 non-null   float64
5   Viscera_weight  4177 non-null   float64
6   Shell_weight    4177 non-null   float64
7   Age             4177 non-null   float64
8   Sex_F           4177 non-null   uint8
9   Sex_I           4177 non-null   uint8
10  Sex_M           4177 non-null   uint8
dtypes: float64(8), uint8(3)
memory usage: 273.4 KB
```

```
In [110... X = x.drop(['Age'], axis = 1)
```

```
In [128... X.head()
```

Out[128]:	Length	Diameter	Height	Whole_weight	Shucked_weight	Viscera_weight	Shell_weight	Sex_F	
0	-0.663474	-0.501673	-1.196422	-0.643390	-0.611770	-0.732343	-0.643590	-0.674834	-
1	-1.601273	-1.572915	-1.330241	-1.259765	-1.219694	-1.236126	-1.257424	-0.674834	-
2	0.006383	0.087510	-0.125873	-0.295144	-0.456142	-0.343709	-0.183214	1.481846	-
3	-0.797445	-0.501673	-0.393511	-0.639118	-0.655541	-0.607596	-0.605225	-0.674834	-
4	-1.779901	-1.680039	-1.597878	-1.303563	-1.268328	-1.322489	-1.372518	-0.674834	

```
In [112... y = x['Age']
```

```
In [129... y.head()
```

```
Out[129]: 0    16.5
1     8.5
2    10.5
3    11.5
4     8.5
Name: Age, dtype: float64
```

9. Scale the independent variables

```
In [114... from sklearn.preprocessing import StandardScaler
```

```
In [115... X_columns = X.select_dtypes(include=np.number).columns.tolist()
X_columns
```

```
Out[115]: ['Length',
           'Diameter',
           'Height',
           'Whole_weight',
           'Shucked_weight',
           'Viscera_weight',
           'Shell_weight',
           'Sex_F',
           'Sex_I',
           'Sex_M']
```

```
In [116... scaler = StandardScaler()
```

```
In [117... X[X_columns] = scaler.fit_transform(X[X_columns])
```

```
In [118... X.head()
```

```
Out[118]:
```

	Length	Diameter	Height	Whole_weight	Shucked_weight	Viscera_weight	Shell_weight	Sex_F
0	-0.663474	-0.501673	-1.196422	-0.643390	-0.611770	-0.732343	-0.643590	-0.674834 -
1	-1.601273	-1.572915	-1.330241	-1.259765	-1.219694	-1.236126	-1.257424	-0.674834 -
2	0.006383	0.087510	-0.125873	-0.295144	-0.456142	-0.343709	-0.183214	1.481846 -
3	-0.797445	-0.501673	-0.393511	-0.639118	-0.655541	-0.607596	-0.605225	-0.674834 -
4	-1.779901	-1.680039	-1.597878	-1.303563	-1.268328	-1.322489	-1.372518	-0.674834

10. Split the data into training and testing

```
In [119... X.shape, y.shape
```

```
Out[119]: ((4177, 10), (4177,))
```

```
In [120... from sklearn.model_selection import train_test_split
x_train, x_test, y_train, y_test = train_test_split(X,y, test_size=0.2, random_state=42)
```

```
In [121... print(' x_train.shape : ',x_train.shape)
print(' y_train.shape : ',y_train.shape)
print(' x_test.shape : ',x_test.shape)
print(' y_test.shape : ',y_test.shape)
```

```
x_train.shape : (3341, 10)
y_train.shape : (3341,)
x_test.shape : (836, 10)
y_test.shape : (836,)
```

Build the Model, Train the Model and Test the Model

```
In [122... #Linear Regression
```

```
from sklearn.linear_model import LinearRegression
```

```
lr = LinearRegression()  
lr.fit(x_train, y_train)  
  
lr_pred = lr.predict(x_test)
```

In [123]...

```
#Random Forest  
from sklearn.ensemble import RandomForestRegressor  
from sklearn.metrics import mean_squared_error, make_scorer  
from sklearn.model_selection import RandomizedSearchCV  
  
rf = RandomForestRegressor()  
  
param = {  
    'max_depth': [3, 6, 9, 12, 15],  
    'n_estimators': [10, 50, 100, 150, 200]  
}  
  
rf_search = RandomizedSearchCV(rf, param_distributions=param, n_iter=5, scoring=make_scorer  
                               n_jobs=-1, cv=5, verbose=3)  
  
rf_search.fit(x_train, y_train)
```

Fitting 5 folds for each of 5 candidates, totalling 25 fits

Out[123]:



In [124]...

```
means = rf_search.cv_results_['mean_test_score']  
params = rf_search.cv_results_['params']  
for mean, param in zip(means, params):  
    print("%f with: %r" % (mean, param))  
    if mean == min(means):  
        print('Best parameters with the minimum Mean Square Error are:', param)  
  
2.640939 with: {'n_estimators': 100, 'max_depth': 6}  
Best parameters with the minimum Mean Square Error are: {'n_estimators': 100, 'max_depth': 6}  
2.700586 with: {'n_estimators': 10, 'max_depth': 6}  
2.652001 with: {'n_estimators': 50, 'max_depth': 6}  
2.908811 with: {'n_estimators': 50, 'max_depth': 3}  
2.744786 with: {'n_estimators': 150, 'max_depth': 15}
```

In [125]...

```
rf = RandomForestRegressor(n_estimators=50, max_depth=6)  
rf.fit(x_train, y_train)  
  
rf_pred = rf.predict(x_test)
```

14. Measure the performance using Metrics

In [137]...

```
from sklearn import metrics  
RMSE1 = np.sqrt(metrics.mean_squared_error(y_test, lr_pred))  
MAE = metrics.mean_absolute_error(y_test, lr_pred)  
MSE = metrics.mean_squared_error(y_test, lr_pred)  
r2_score(y_test, lr_pred)
```



```

print('Linear Regression :')
print('----- ')
print('MAE:', MAE)
print('MSE:', MSE)
print('RMSE:', RMSE1)
print('R2 Score :', R2)
print('\n\n')

```

```

Linear Regression :
-----
MAE: 1.3252222883409852
MSE: 2.9793407734555917
RMSE: 1.726076699760353
R2 Score : 0.4480445148439044

```

```

In [138... from sklearn import metrics
RMSE2 = np.sqrt(metrics.mean_squared_error(y_test, rf_pred))
MAE = metrics.mean_absolute_error(y_test, rf_pred)
MSE = metrics.mean_squared_error(y_test, rf_pred)
R2 = metrics.r2_score(y_test, rf_pred)
print('Random Forest Contains:')
print('----- ')
print('MAE:', MAE)
print('MSE:', MSE)
print('RMSE:', RMSE2)
print('R2 Score :', R2)

```

```

Random Forest Contains:
-----
MAE: 1.246195573419611
MSE: 2.532955077408326
RMSE: 1.5915260215932148
R2 Score : 0.5307423504267539

```

Compare Linear Regression and Random Forest

Random Forest got low rmse value than Linear Regression

```

In [139... RMSE = RMSE1-RMSE2
print(RMSE)

```

```

0.13455067816713817

```